

Fast Concurrent Data-Structures Through Explicit Timestamping

Mike Dodds, Andreas Haas, Christoph M. Kirsch
mike.dodds@york.ac.uk, ahaas@cs.uni-salzburg.at, ck@cs.uni-salzburg.at



Abstract

Concurrent data-structures, such as stacks, queues and dequeues, often implicitly enforce a total order over elements with their underlying memory layout. However, linearizability only requires that elements are ordered if the inserting methods ran sequentially. We propose a new data-structure design which uses explicit timestamping to avoid unwanted ordering. Elements can be left unordered by associating them with unordered timestamps if their insert operations ran concurrently. In our approach, more concurrency translates into less ordering, and thus less-contended removal and ultimately higher performance and scalability.

Key Ideas

- Order elements in the data-structure only partially by using explicit timestamping.
- Store elements in thread-local buffers to avoid synchronization of insert operations.
- Use the RDTSCP CPU instruction for highly-scalable timestamp generation.
- Use intervals as timestamps to reduce the order on timestamps while still providing linearizability.

TS Deque Pseudo Code

```
TS_Deque {
    TS_Buffer buffer;
    void insertR(Element element) {
        item = buffer.insR(element);
        timestamp = buffer.newTimestamp();
        buffer.setTimestamp(item, timestamp);
    }
    Element removeR() {
        do {
            item = buffer.tryRemoveR();
        } while (!item.isValid());
        if (item.isEmpty()) {
            return EMPTY;
        } else {
            return item.element;
        }
    }
}
```

insertL and **removeL** are defined analogously.

Correctness

- The TS deque is linearizable with respect to the sequential specification of a deque.
- The remove operations of the TS deque are lock-free, the insert operations are wait-free.

Additional Informantion

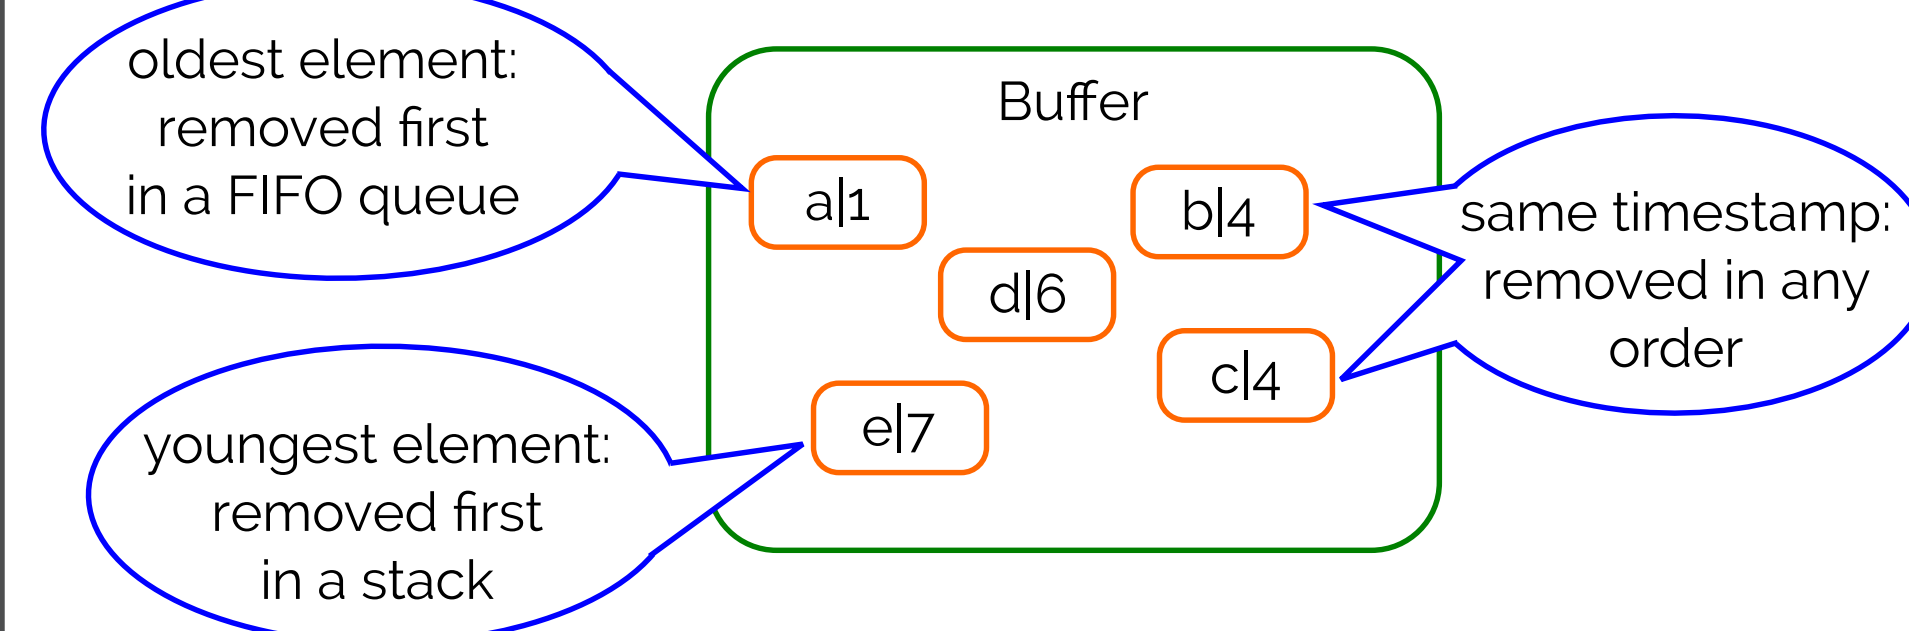
<http://scal.cs.uni-salzburg.at/tsdeque>

Acknowledgements

This work has been supported by the National Research Network RiSE on Rigorous Systems Engineering (Austrian Science Fund S11404-N23).

Partially Ordered Elements

Elements with timestamps are stored in an unordered buffer. Elements which were inserted concurrently may have the same timestamp.

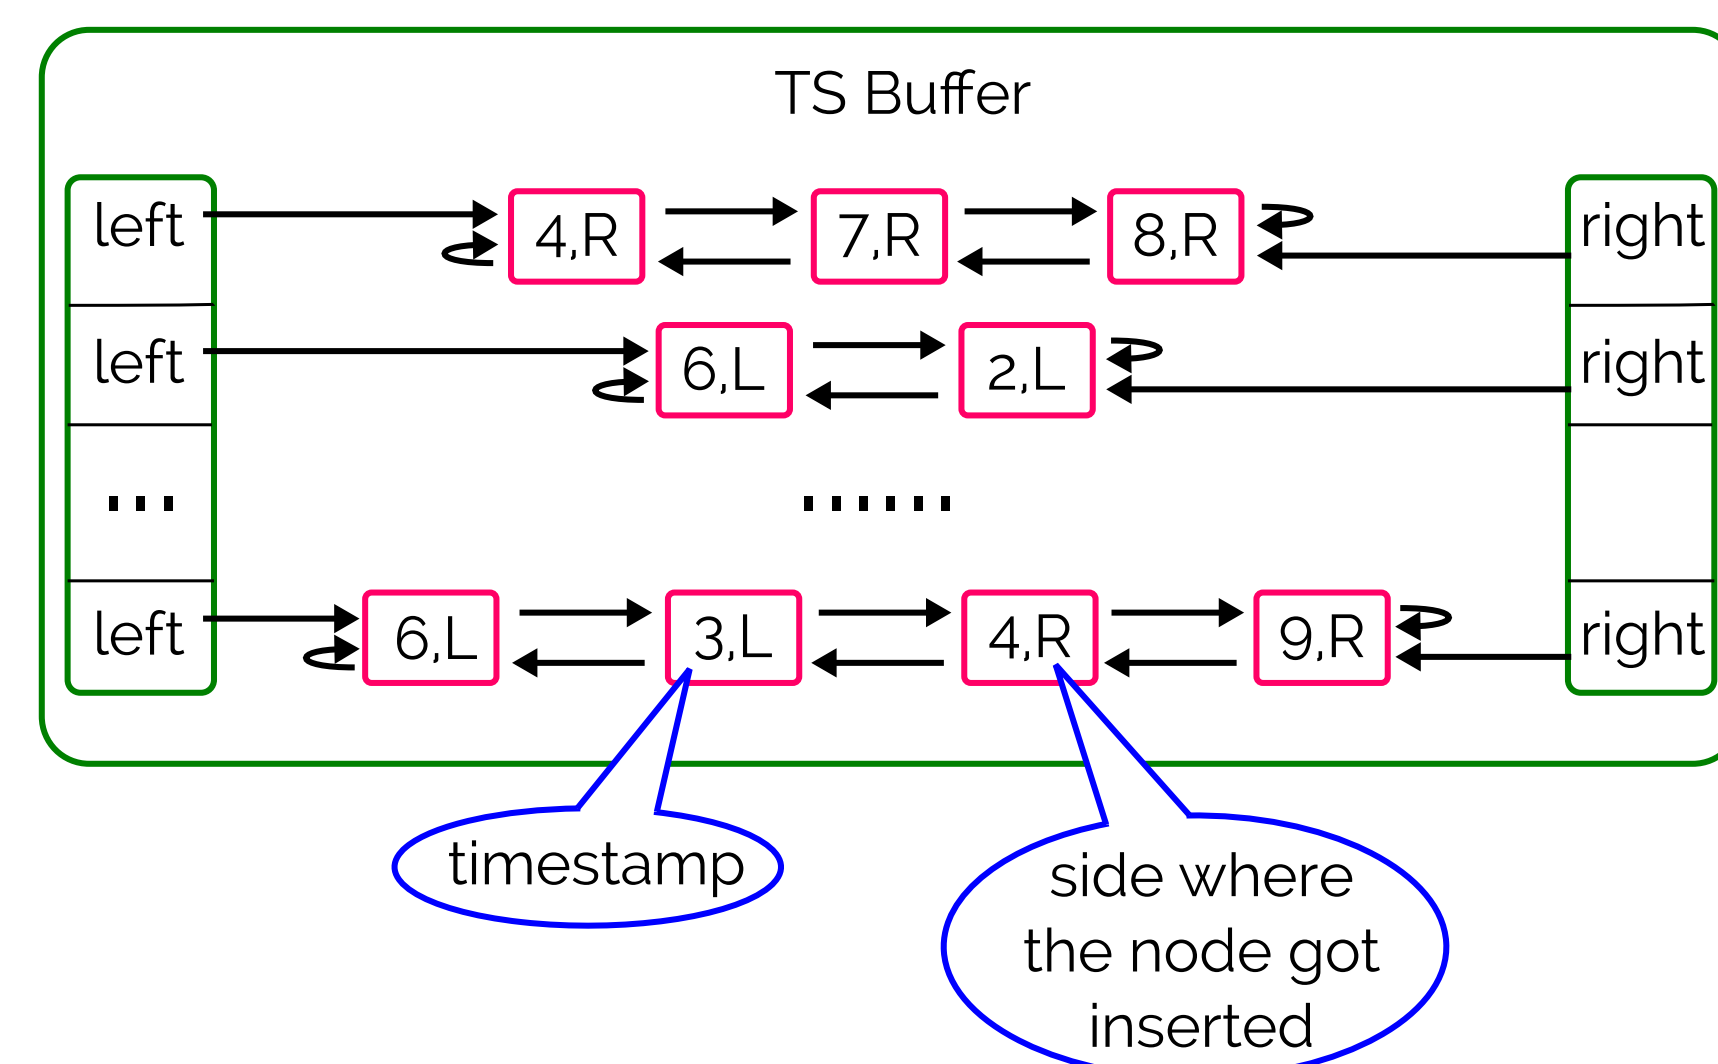


TS-Buffer Design

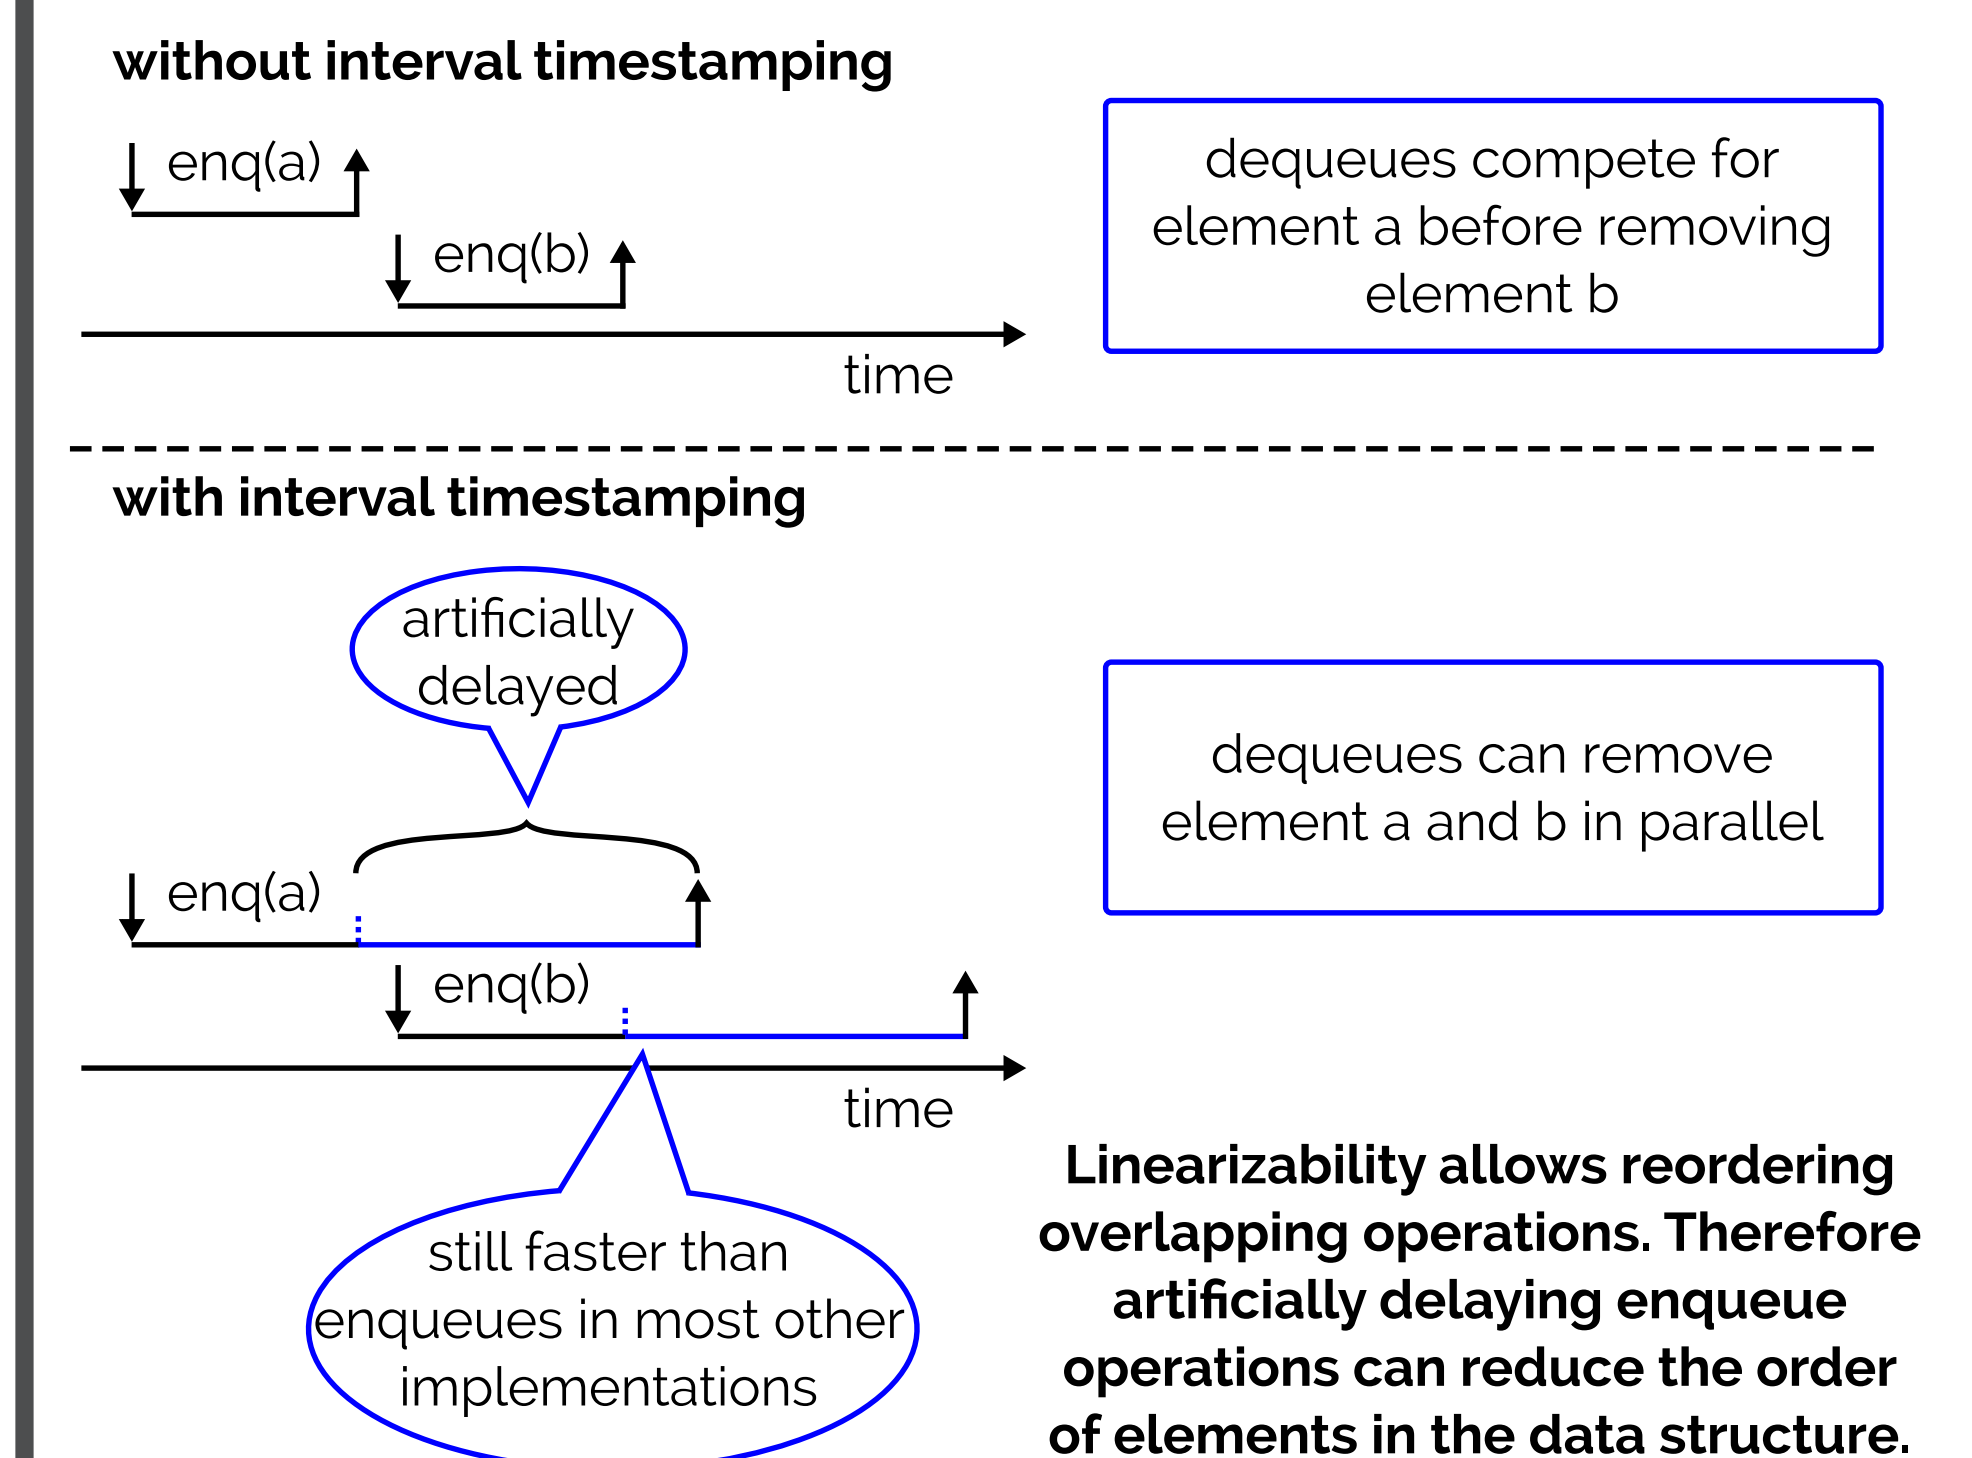
insR: Each thread inserts at the right side of its own linked list.

tryRemoveR: A thread searches through the right ends of all linked lists for the right-most element to remove. Nodes are removed logically first by setting a **removed** flag. Unlinking occurs in later insert or remove operations. If setting the **removed** flag fails, then **tryRemoveR** returns an invalid item.

insL/tryRemoveL: Analogous to **insR** and **tryRemoveR**, but the left side of the linked lists is accessed.




Interval Timestamping



Interval Length

Performance of the TS deque with interval timestamping with an increasing delay in a high-contention producer-consumer benchmark. All measurements are done with 80 threads on a 40-core (2 hyperthreads per core) server machine. The left y-axis shows the average number of operations (both insert and remove operations) per millisecond, the right y-axis shows the average number of CAS retries per remove operation.



Performance in a Producer/Consumer Benchmark

Performance and scalability of the TS deque in a high-contention producer-consumer benchmark with an increasing number of threads. The baselines are a Treiber stack, an elimination-backoff stack (EB stack), a Michael-Scott queue (MS queue), and a flat-combining queue (FC queue). Stack-specific and queue-specific TS-Buffers (Hardcoded TS Stack and Hardcoded TS Queue) allow additional performance gains.

